

Sizing Matters

gappy

September 15, 2026

How large should a position be? We can usually explain why we like a stock more easily than why we own twice as much of it as another one. Yet, every portfolio holds positions that twice as big as the smallest ones in the same portfolio. Sometimes more than one position, sometimes more than twice as big. These sizing decisions are material to your performance as a portfolio manager or as an analyst. In this note we analyze the problem and we make some practical recommendations. The material is organized into three sections. First, we look at the evidence: how much sizing matters on realized performance of investors. We use anonymized BAM data. Secondly, we analyze the sizing problem analytically. We close with a list of practical recommendations.

1 What the Data Say

Let's give away the message first. Selection makes the larger contribution among profitable PMs, while sizing contributes enough to matter materially to their results. About half of all PMs have no sizing skill whatsoever, even though all of them believe their are good at it. The goal of this report is to bring PnL dreams and reality closer together.

To compare the two sources of PnL, we need to define them and measure them. Consider a portfolio on a given day. We decompose it into two portfolios:

- A portfolio that preserves the direction of each position and gives every position the same absolute Gross Market Value, equal to the average absolute GMV of the original positions. This is the *selection* portfolio.
- A portfolio whose positions are the differences between the original positions and the equally sized positions. This is the *sizing* portfolio.

The two portfolios sum to the original portfolio. We compare their cumulative, launch-to-date (LTD) PnLs using a filtered sample of 155 ACUs active at some point since 2021, each with more than 150 observations and an assigned theme. Each ACU counts as one PM observation in the tables. Profitability is determined by reported total PnL. In the supplied summary data, the reported selection and sizing PnLs do not generally sum to that total; we therefore distinguish shares of their combined PnL from ratios to reported total PnL.

Table 1 shows how often each component contributes positively. Among the 80 profitable PMs, selection PnL is positive for 89%, compared with 72% for sizing. Selection PnL exceeds sizing PnL for 68% of these PMs. Selection is thus both more frequently positive and more often the larger contributor. Yet sizing is positive for nearly three quarters of profitable PMs, and the median ratio of sizing PnL to reported total PnL is 25%. A contribution of that magnitude deserves attention.

Table 1: Selection and sizing outcomes by reported total PnL, by PM. Percentages are computed within each group. The median ratio is sizing PnL divided by reported total PnL and is shown only for profitable PMs. The nonpositive group includes six zero-PnL observations.

Total PnL	% (Sel. > Siz.)	% Siz.PnL > 0	% Sel.PnL > 0	median(% Siz.PnL/Tot.PnL)	# PMs
nonpositive	45	35	41		75
positive	68	72	89	25	80

The other 75 PMs have nonpositive total PnL, including six with zero total PnL. In this group, sizing PnL is positive for 35% and selection PnL for 41%. Positive sizing can coexist with an overall loss, just as positive selection can. The comparison favors selection without implying that either component alone determines profitability.

Table 2 measures the relative dollar contributions among profitable PMs. We first sum each component across PMs within a theme, then divide by the sum of selection and sizing PnL. Across all 80 profitable PMs, selection contributes 76.9% of the combined PnL and sizing contributes 23.1%: selection contributes about 3.3 times as much. This is the strongest direct evidence for the relative importance of the two decisions. Sizing supplies almost a quarter of their combined profits, while selection supplies just over three quarters.

Table 2: Shares of combined selection and sizing LTD PnL among the 80 profitable PMs. Each component is summed across PMs within a theme and divided by the sum of both components. The ALL row pools all profitable PMs.

Theme	Aggregate % Selection PnL	Aggregate % Sizing PnL	# PMs
CONSUMER	79.3	20.7	9
FINANCIAL & REIT	52.6	47.4	9
GENERALIST	65.0	35.0	11
HEALTHCARE	52.8	47.2	9
IND & ENE & UTILS	88.4	11.6	21
TMT	78.2	21.8	21
ALL	76.9	23.1	80

Selection has the larger aggregate share in every theme, but the gap varies substantially. Sizing contributes 11.6% in the combined industrial, energy, and utilities group, about 21% in consumer and TMT, and 35.0% among generalists. In financials and REITs, and in healthcare, sizing contributes approximately 47%, close to selection's contribution. The overall ordering therefore allows an important role for sizing in particular groups. These shares describe the profitable subset; the smaller groups contain only nine PM observations each.

Table 3 returns to all 155 PM observations and compares their LTD outcomes. Selection PnL is more strongly correlated with total PnL than sizing PnL is: 0.80 versus 0.61. The same ordering holds for alpha PnL, with correlations of 0.90 and 0.67. Sizing remains meaningfully associated with overall results. Its correlation with selection is only 0.34, so strong selection outcomes do not automatically come with strong sizing outcomes. These are correlations across PMs' cumulative PnLs; they measure neither shares of profits nor diversification through time.




Figure 1: Selection and sizing LTD PnL by theme and overall profitability. Each point is one PM observation; the blue line marks equal PnL.

Table 3: Cross-sectional correlations of LTD PnL across all 155 PM observations, including profitable and nonprofitable PMs. The entries are correlations of cumulative outcomes, not daily PnL correlations.

	Total PnL (\$mm)	Alpha PnL (\$mm)	Selection PnL (\$mm)	Sizing PnL (\$mm)
Total PnL (\$mm)	1.00	0.91	0.80	0.61
Alpha PnL (\$mm)	0.91	1.00	0.90	0.67
Selection PnL (\$mm)	0.80	0.90	1.00	0.34
Sizing PnL (\$mm)	0.61	0.67	0.34	1.00

Figure 1 makes the variation across PMs visible: points below the blue line have larger selection than sizing PnL, while points above it have larger sizing PnL. Notice how the contribution to sizing for Healthcare and Financial and REIT is driven by two points in each case. The results make sizing a material part of the investment process while giving selection the larger role. They describe realized outcomes and do not, by themselves, establish how much a particular change in sizing would improve future performance.

Insight 1: The contribution of sizing

Sizing contributes approximately 25% of PnL to a fundamental PM strategy.

2 Thinking about sizing

2.1 Start with the right chapter and the right inputs

Seen from 10,000 meters, the sizing process might look like Figure 2: risk, expected return, and market impact all appear to feed directly into the position size.




Figure 2: A bird's eye view of positions. The yellow boxes collect the inputs, and the green box represents the resulting position size.

But this picture contains both too much and too little. It contains too much, because the market impact model *should not* be part of the decision process of the portfolio manager. Instead, it should be part of the *execution decision process* which receives as an input the position size *and* the market impact model in order to decide the time and size of the trade orders. Figure 3 separates these two stages.




Figure 3: A more detailed picture. Risk and expected return feed the green position-size decision. The grey boxes show execution: the target position and the market impact model determine the trade orders.

The picture's resolution is too low: an expected return alone does not describe the full investment forecast. We need the uncertainty attached to that forecast and the horizon over which the return is expected to be earned. Figure 4 makes these three inputs explicit alongside the risk model.




Figure 4: Return and risk models. The return model groups expected return, return uncertainty, and return horizon. The risk model groups idiosyncratic risk and factor risk. Both supply inputs to the position-size decision.

Now, we discuss how to include these four variables in the sizing decision process.

1. Return model.

- (a) **Expected return.** This is the main input and is obviously the most important. However, *how* to convert the output of fundamental research into a single number is not obvious. We will revisit this subject in Section 3.
- (b) **Return uncertainty.** Uncertainty of the forecasts (as opposed to volatility of stock returns) is also intuitively important. *Uncertainty is a way to quantify conviction.* Stating “I have high conviction in a stock” is ambiguous. It may mean “I believe the stock will perform better than its peers”, but also “my belief about *this belief* is accurate”. Separate the two. Return uncertainty describes the error that you believe you may make. We estimate expected returns. Correspondingly, we estimate standard errors of these

forecasts as a way to quantify uncertainty. *We recommend thinking about the standard error as a percentage of the expectation.*

- (c) **Return horizon.** It is important that all returns are expressed over the same horizon. Otherwise we are comparing incommensurable quantities: a stock return forecast for the earnings next week is not the same as a one-year view on another stock. A possible 10% gain in three months and one in three years are different opportunities. How to use the error is a problem that deserves its own section. For the time being, we address the issue of multiple horizons. An effective process is to *group stocks by forecast horizons, and then adjust all of them to the same horizon*. The process is simple and we first illustrate it with an example. Say that you forecast the return of a stock to be 20% one year out, and assess a standard error of 30%. All your other forecasts are a quarter out (or 63 trading days). You shift the forecast horizon of the expected return by dividing the expected return by 4. The formula is

$$\begin{aligned} (\text{quarterly expected return}) &= (\text{annualized expected return}) \times \frac{(\text{one quarter})}{(\text{one year})} \\ &= (\text{annualized expected return}) \times \frac{1}{4} \end{aligned}$$

What about the error? The same way that we convert volatility of a portfolio over different horizons by doing a “square root of time” multiplication, the same is the case for error:

$$\begin{aligned} (\text{quarterly std.error of return}) &= (\text{annualized std.error of return}) \times \sqrt{\frac{(\text{one quarter})}{(\text{one year})}} \\ &= (\text{annualized std.error of return}) \times \frac{1}{2} \end{aligned}$$

The insight here is that if you make a forecast at horizons (one year, or two), and you convert it to a shorter horizon, its signal-to-noise worsens. Since, intuitively, more uncertainty implies less confidence in the forecast and underweight, this means that *you should penalize it more, compared to other genuinely short-term forecasts*. This is very intuitive, and yet it is often the case that investors only rescale expectations, but not errors.




Figure 5: Selection and sizing LTD PnL by theme and overall profitability. Each point is one PM observation; the blue line marks equal PnL.

2. Risk model.

- (a) **Factor risk.** The central idea of factor modeling is to *separate sources of returns that you can forecast from those that you can't*. How much of the expected return comes from the company, and how much from the market, its industry, or a style such as value? You don't want to be too strict in controlling factor risk, but *your goal is always to forecast company-specific, idiosyncratic returns*. Use the easy-to-use simple optimizer in RD2. If you can't get to a solution that you like, move to the more sophisticated one. Importantly, *don't bet on factors*. Maybe, if you have a strong conviction on a specific factor to have a future positive return, you can require that exposure in the factor be non-negative. In general, time spent guessing factor returns is time subtracted from what you do best: fundamental research.
- (b) **Idiosyncratic risk.** Mean-variance optimization prescribes a simple formula for sizing when company-specific returns are uncorrelated:

$$(\text{dollar position}) = (\text{portfolio scale}) \times \frac{(\text{expected return})}{(\text{return volatility})^2}.$$

Here, expected return and return volatility refer to the company-specific component, after removing factor exposure. The portfolio scale is a common positive multiplier chosen to fit the portfolio's risk budget. A positive position is a long; a negative position is a short. All return and risk inputs must refer to the same holding period.

The formula is extremely sensitive to volatility. Suppose we forecast the same return for two stocks, one with annualized idiosyncratic volatility of 15%, the other with 45%. The second stock is three times as volatile, but its variance is nine times as large. It therefore receives one ninth of the first stock's position.

The problem is that *the same denominator that rewards low volatility also magnifies errors in expected returns*. If we are wrong about the quieter stock, we are wrong in a

much larger position. At a fixed portfolio scale, the same error in its return forecast changes its recommended position nine times as much. Low return volatility does not, by itself, mean that our estimate of expected return is precise.

One response is to divide the forecast by volatility instead of variance:

$$(\text{dollar position}) = (\text{portfolio scale}) \times \frac{(\text{expected return})}{(\text{return volatility})}.$$

The position ratio in our example falls from nine to three. Stocks with the same absolute expected return then receive the same standalone dollar volatility: the absolute dollar position multiplied by return volatility. This is the sense in which the rule is called *risk parity sizing*. It is more conservative about concentrating dollars in low-volatility stocks, although it need not produce a portfolio with lower total risk. With different forecast magnitudes, dollar volatility is proportional to the magnitude of the expected return, rather than equal across stocks.

The main drawback is that the rule has no general optimality justification. It can perform better in practice, but that result depends on the forecasts and the investment universe. Replacing variance with volatility also leaves out an input we have already asked the analyst to quantify: *uncertainty about the expected return*.

Measure that uncertainty by the standard error of the expected-return estimate. If it is independent of the subsequent return fluctuations around the unknown expected return, the two sources of variance add. Mean-variance sizing then becomes¹

$$(\text{dollar position}) = (\text{portfolio scale}) \times \frac{(\text{expected return})}{(\text{return idio volatility})^2 + (\text{forecast standard error})^2}.$$

Return volatility measures fluctuations around the unknown expected return; the forecast standard error measures uncertainty about that expected return itself. We should not add forecast uncertainty a second time if it is already included in the risk estimate. This stock-by-stock formula also assumes that neither source of risk is correlated across stocks. Shared business risks or common errors in our forecasts require us to account for how positions move together.

Forecast uncertainty can be large enough to dominate idiosyncratic variance. When it accounts for most of the combined variance, the formula simplifies to

$$(\text{dollar position}) \approx (\text{portfolio scale}) \times \frac{(\text{expected return})}{(\text{forecast standard error})^2}.$$

If forecast standard errors are also roughly the same across stocks, the denominator is approximately common: *relative position sizes should then follow expected returns*. Merely having a standard error greater than volatility is not enough for this approximation; forecast uncertainty must account for most of the combined variance.

Our practical recommendation is to begin with a roughly common forecast standard error unless we have evidence that some forecasts are more precise than others. When this uncertainty dominates return risk, proportional sizing is a reasonable first approximation. A stock with twice the expected return receives twice the position. Better evidence about forecast precision should refine that allocation through the full denominator, while the common scale remains a separate portfolio risk decision.

¹A separate appendix on uncertainty-adjusted portfolios derives this result.

Table 4 collects these rules in compact notation. Each rule needs its own portfolio scale. Because the rules use risk differently, applying the same numerical multiplier would not give a meaningful comparison.

Table 4: Sizing rules and the information they use. For stock i , x_i is the signed dollar position, α_i the expected idiosyncratic return, σ_i the idiosyncratic return volatility, and τ_i the standard error of the expected-return forecast. The positive scaling factor κ is common across stocks within a rule and is calibrated separately for each rule. All inputs use the same holding period.

Rule	Dollar position x_i	Interpretation
Mean-variance	$\kappa \frac{\alpha_i}{\sigma_i^2}$	Uses expected return and idiosyncratic variance; treats the expected return as known.
Alpha over volatility	$\kappa \frac{\alpha_i}{\sigma_i}$	Equal standalone dollar volatility for equal absolute forecasts.
Uncertainty-adjusted	$\kappa \frac{\alpha_i}{\sigma_i^2 + \tau_i^2}$	Includes uncertainty about expected returns in the variance penalty.
Uncertainty dominates	$\kappa \frac{\alpha_i}{\tau_i^2}$	Approximation when $\tau_i^2 \gg \sigma_i^2$.
Proportional	$\kappa \alpha_i$	Follows when dominant forecast uncertainty is approximately common across stocks.

Proportional sizing gives equal absolute dollar positions only when all buy and sell signals have the same absolute strength. With numerical return forecasts, it preserves their relative magnitudes.

Table 5 shows the effect on a concrete position. Consider a stock with a 10% expected idiosyncratic return over one year. For each rule, choose the scale so that the position is \$10 million when volatility is 30% and forecast uncertainty is zero. Then hold that rule's scale fixed as the inputs change. Each row is an alternative scenario for the same stock.

Table 5: Position size under different volatility and forecast-uncertainty assumptions. The one-year return forecast is 10% throughout; positions are in millions of dollars.

Idiosyncratic volatility	Forecast standard error	Mean-variance	Alpha over volatility	Uncertainty-adjusted
15%	0%	40.00	20.00	40.00
30%	0%	10.00	10.00	10.00
45%	0%	4.44	6.67	4.44
15%	30%	40.00	20.00	8.00
30%	30%	10.00	10.00	5.00
45%	30%	4.44	6.67	3.08
15%	90%	40.00	20.00	1.08
30%	90%	10.00	10.00	1.00
45%	90%	4.44	6.67	0.89

Without forecast uncertainty, mean-variance sizing assigns nine times as much to the 15%-

volatility stock as to the 45%-volatility stock; alpha-over-volatility sizing assigns three times as much. With a common 90% forecast standard error, the uncertainty-adjusted positions are much closer: \$1.08 million and \$0.89 million, a ratio of about 1.22. The deliberately large 90% assumption illustrates how positions converge as forecast uncertainty increases. The proportional rule, under its reference calibration, would assign \$10 million in every row because the expected return never changes.

The uncertainty adjustment therefore changes both relative sizing and overall exposure. In practice, you want to hold a portfolio that has a total volatility compatible with your risk appetite and with the budget you have received from the firm. You can achieve this goal by multiplying all the positions by the same constant.

3 Estimating Returns and Their Uncertainty

Estimating a stock's expected return and the uncertainty around that estimate involves subjective judgment. The analyst must decide which evidence to trust, how the business will develop, what the market already expects, and when new information might change the share price. Thoughtful analysts can examine the same company and reach different conclusions about both the likely payoff and their confidence in it.

Subjective does not mean unstructured. Best practices and checklists can make the assumptions explicit, expose gaps in the research, and help a team apply consistent standards. They do not turn a conviction score into a calibrated probability or a numerical standard error. Their purpose is to explain why a forecast deserves more or less confidence, and what evidence would change it.

The following checklist covers the research behind a forecast, the assessment of its uncertainty, the translation of that judgment into a position, and the process for recording, updating, and reviewing decisions. The examples illustrate how to apply each recommendation. The labels indicate whether a recommendation concerns estimating the payoff or its timing (models expectation), assessing evidence, forecast reliability, or adverse outcomes (models uncertainty), or both. Operational recommendations are labelled (workflow).

1. **Record forecasts, reasoning, and open research questions in a shared format.** (workflow)
Use a common data format and schema for scenarios, price targets, expected returns, forecast uncertainty, qualitative assessments, and supporting notes. Record the forecast date and horizon, and link the inputs to the supporting financial model so that the stored view and model remain consistent. Capture new research angles when they arise, including non-obvious meeting findings and the work needed to test them.
Example: a stock's record could contain its three-month forecast, upside and downside prices, forecast standard error, and model version, alongside a question about why service revenue differs from the installed equipment base and a task to check contract coverage and renewal pricing.
2. **Quantify the fundamental disagreement with market expectations.** (models expectation)
State the reporting period, what the company will earn, how far that earnings or cash-flow estimate lies above or below prevailing forecasts, and how the difference could affect the share price. Build the case around operating and accounting drivers that can be investigated. Give more weight to a material, well-supported disagreement than to a small difference that

normal forecasting error could erase. Consider both the likelihood and size of the payoff: an earnings beat may already be priced in or offset by a lower valuation multiple.

Example: a forecast of \$1.60 in earnings per share against a highest broker estimate of \$1.10 presents a larger disagreement than a forecast of \$1.12, but the additional earnings must be explained by a measurable driver, such as improved product profitability.

3. **Form an independent view before relying on other investors' opinions.** (models expectation) Build the forecast from company reports, presentations, accounts, and operating evidence. Use prevailing forecasts to understand the disagreement, without borrowing their assumptions unexamined. Build the case around the expected fundamental development; avoid making its success depend on persuading a broker to publish a supportive note.
Example: model an expiring order from the company's disclosures before checking whether consensus has mistakenly carried that revenue into the following year.
4. **Protect research time from routine work and market noise.** (workflow) Automate or delegate recurring tasks and limit distractions that do not change the investment case. Constantly monitoring small price moves, messages, and news can crowd out the work needed to find a material opportunity.
Example: automate the weekly update of consensus estimates and use the time saved to investigate a revenue line that does not reconcile with reported shipment volumes.
5. **Build a cumulative knowledge base and seek situations where that knowledge matters.** (models expectation) Develop a deep understanding of the industry's economics, accounting, products, and competitive relationships. Repeated research makes it possible to connect facts that appear unrelated to a less experienced observer. Seek complex, company-specific situations where that accumulated knowledge provides an advantage.
Example: earlier work on an engine manufacturer can help establish how an aircraft producer's production cut would affect the supplier.
6. **Build a detailed operating model and connect outside evidence to it.** (expectation, uncertainty) Model one level below the company's disclosure, breaking reported divisions into the drivers of revenue, profit, and cash flow. Understand revenue recognition, cost timing, and working capital. Before using customer observations, surveys, competitor commentary, or third-party data, establish what they measure and how they map to those drivers. Check seasonality, comparison periods, one-off items, and when the effect enters the accounts.
Example: model airline passenger volumes and yields separately, and map booking data to the relevant travel period before estimating revenue. Strong bookings should not be inserted directly as growth in this quarter's reported revenue.
7. **Find what others have missed, then quantify whether it matters.** (models expectation) Knowing the public information is the starting point. Look for an overlooked fact or an inconsistent assumption, explain its financial consequence, and estimate its importance before committing extensive research time. Explain why the market may have missed it and what would happen if it does not materialize. A novel detail that barely changes earnings is not necessarily a useful investment idea.
Example: if forecasts incorporate a stronger dollar's effect on airline costs but omit the offsetting dollar revenues, reconciling both sides of the currency exposure can reveal a material forecasting error.
8. **Prepare management and industry conversations around unanswered questions.** (expectation, uncertainty) Go beyond the current quarter's trading update. Ask about the mechanisms that could change earnings over the next several quarters or years, and use

customers, suppliers, partners, and experts to examine the same issues from different perspectives. Prepare the agenda and follow up on inconsistencies.

Example: ask how a disclosed backlog converts into revenue over the next two years, what could delay delivery, and whether suppliers' capacity plans support management's schedule.

9. **Assess management competence and transparency separately.** (models uncertainty) A capable team can withhold bad news, while a candid team may have poor internal information. Evaluate the quality of management's forecasting and controls separately from its willingness to explain developments honestly. Reconcile its statements with the accounts and operating evidence.

Example: a persuasive CFO who repeatedly postpones disclosure of cost overruns should not receive the same transparency assessment as one who flags the problem early, even if both understand the business well.

10. **Separate dependable drivers from informed guesses.** (models uncertainty) Assess the source of each part of the earnings difference. An identifiable accounting effect or an expiring contract may be easier to establish than next quarter's demand, price, or volume. Several similar data points can reinforce the same uncertain assumption without making it much more reliable.

Example: an order expiration specified in a contract provides firmer evidence than a 9% organic-growth forecast against consensus of 6%. If growth instead comes in at 4%, the apparent forecasting advantage disappears.

11. **Investigate and stress the drivers that matter most to the forecast.** (models uncertainty) Direct the strongest research effort toward the assumptions that account for most of the earnings or cash-flow variation. Ask how much of the disagreement survives if an important driver disappoints. A large forecast difference is useful only if it is not entirely dependent on one fragile assumption.

Example: if an illustrative \$1.60 earnings forecast falls to \$1.35 when a planned price increase is removed, it still exceeds a \$1.10 consensus forecast. That exercise identifies both the remaining opportunity and the assumption that needs monitoring.

12. **Identify the exposures you cannot forecast and consider hedging them.** (models uncertainty) Concentrate the investment case on the drivers the research supports. Identify other exposures that could dominate the outcome and consider whether they can be reduced or hedged. Choose hedges for identifiable exposures and assess the risks they leave behind; not every unwanted exposure can be eliminated.

Example: an airline thesis based on improved passenger mix may still carry substantial fuel-price exposure. Assess a feasible hedge and the remaining mismatch before treating the position as a pure passenger-mix bet.

13. **Estimate and keep updating the remaining upside and downside.** (expectation, uncertainty) Translate operating forecasts into possible share prices using the valuation measures the market applies, such as earnings, sales, EBITDA, cash flow, or discounted cash flow. Compare the team's and the street's estimates with history and peers, and include realistic price moves around the next catalysts. Revisit the payoff after price changes or new information; holding or taking profits should depend on what remains, not on gains already earned.

Example: a healthcare company's attractive long-term cash-flow valuation must still allow for a sharp decline after an unsuccessful trial. Conversely, a completed price rise need not

exhaust the upside: with free cash flow unchanged, moving from three times to twice a reference cash-flow yield requires a further 50% price increase. This calculation measures a possible move, not its likelihood.

14. **Map the catalyst sequence and the evidence that would invalidate the thesis.** (expectation, uncertainty) Identify what could cause the market to revise its view, when it could happen, and what comes after the first event. Distinguish a delayed catalyst from evidence that the investment case is wrong. Consider whether management's explanation or a later event could reverse the initial price reaction.

Example: a thesis may depend on earnings, then an investor day, then a product launch. If the earnings release is inconclusive, specify what the next two events must show to justify continuing to hold the position.

15. **Compare the near-term catalyst with the longer-term thesis before choosing size and holding period.** (expectation, uncertainty) Write down both the next-event forecast and the earnings outlook over several years; ask whether you would still own or short the company without the imminent event. When the horizons conflict, adjust size, timing, or holding period: a favorable event in a weak business may justify a brief, liquid trade, while a strong business may warrant a smaller position before an adverse event. When the horizons align, consider a larger allocation only if independent research supports a material earnings difference, dependable drivers, and substantial remaining payoff. Discuss concentration with the PM and risk team, and ensure the portfolio can tolerate and finance the path to resolution.

Example: a shipment pulled forward may justify only a one-week trade in a structurally weak company, with a planned exit. An imminent backlog update that supports both a near-term earnings revision and multi-year growth can justify discussing a larger, longer-held position.

16. **Give every idea an explicit dollar loss budget.** (models uncertainty) State how much the portfolio can afford to lose on the position, including ideas intended only as trades. Combine that budget with the adverse price move in the scenario, and consider the path around the event rather than only the final target price.

Example: a \$500,000 loss budget and a 25% adverse-move scenario imply a \$2 million position before other constraints. The scenario guides the allocation; it does not guarantee that losses stop at the budget.

17. **Use a shared checklist for qualitative assessments.** (models uncertainty) Write down the questions used to assess fundamentals, the financial model, and market expectations, and share the checklist with analysts. Define what strong and weak evidence look like so that assessments do not depend only on the person's mood or presentation. A one-to-three score for each dimension is one possible convention; keep these research assessments distinct from any portfolio sizing rank.

Example: an idea could receive a three for the model but a one for the setup because the team has not established what investors already expect. The checklist makes that missing work explicit.

18. **Set dollar targets across the book and review deviations from them.** (workflow) Define the portfolio ranks; in a one-to-five system, one might mean a starter position or brief trade and five the largest intended category. Document how ranks, stock volatility, the portfolio budget, competing ranks, and the extra weight given to the best ideas produce dollar targets. Equally ranked ideas can need different exposures, and new high-ranked ideas compete for

the same budget. Compare actual holdings with targets to identify underfunded ideas and excessive risk. The PM remains responsible for increases, reductions, and deliberate, documented overrides.

Example: if stronger earnings evidence raises a medtech short's model target to \$60 million from an actual \$20 million holding, adding \$10–20 million can narrow the gap while retaining discretion over the final exposure.

19. **Test the current size by considering an adverse price move.** (models uncertainty) Ask whether the PM would add to a long after a decline or to a short after a rise. Hesitation while already holding a large position can be a warning that the size exceeds conviction. Apply the check only after reconsidering the thesis and the remaining loss budget; willingness to add is not itself evidence of a good forecast.

Example: after a long falls 5%, distinguish an unchanged earnings case at a better price from a decline caused by a lost customer that invalidates the forecast.

20. **Refresh the inputs at least as often as the typical investment horizon.** (expectation, uncertainty) Use performance attribution to examine the holding period at which the alpha curve peaks, and compare it with profitability by idea age. These measures should broadly inform the same review horizon; investigate a material difference. Review forecasts at least at that interval, whether it is several weeks or a quarter, and revisit them sooner when material information arrives.

Example: if alpha tends to peak after about 60 trading days and older ideas become less profitable, review each thesis and its sizing inputs at least quarterly rather than carrying an unchanged forecast indefinitely.

21. **Use every earnings release to improve the forecast and its confidence assessment.** (expectation, uncertainty) Reconcile outcomes with the model, identify the assumptions that failed, and update the research process. Treat curiosity, persistence, and humility as part of a repeated cycle of learning. Over time, the record of forecast errors can make the next conviction assessment better informed.

Example: if revenue misses because a large order was recognized later than modeled, revise the recognition schedule and the questions asked of management. Distinguish that timing error from a lost order that changes the longer-term earnings case.

22. **Review sizing outcomes with the team every quarter and improve the checklist.** (workflow) Examine selection and sizing PnL, then revisit the position sizes of the biggest winners and losers. Compare the decisions with the forecasts, uncertainty assessments, and constraints recorded at the time. Discuss whether the outcome exposed a weak assumption, an inconsistent decision, or a missing case, and update the shared checklist accordingly.

Example: if several large losses came from treating uncertain demand forecasts as established facts, add an explicit demand-scenario check and require the team to explain how that uncertainty affects size.

Procedure 1: A Sizing Checklist

1. **Group stocks by forecast horizon.** Separate ideas according to when you expect the forecast return to be earned.
2. **Estimate returns and assign uncertainty within each group.** Use your own process and, where useful, the best practices in Section 3. Express each forecast's standard error as a percentage of its absolute expected return. For example, uncertainty of 50% on a forecast of +10% or −10% means a standard error of 5 percentage points.
3. **Convert all inputs to one common horizon.** Using the time-scaling convention described above, scale expected returns in proportion to time and forecast standard errors with the square root of time. Express idiosyncratic return volatilities over that same horizon.
4. **Compute the raw position sizes.** Use the uncertainty-adjusted rule:

$$(\text{raw size}) = \frac{(\text{expected return})}{(\text{idiosyncratic volatility})^2 + (\text{forecast standard error})^2}.$$

Use company-specific expected returns and consistent units for all inputs. Positive sizes are longs; negative sizes are shorts. These are relative sizes before setting the portfolio's overall risk.

5. **Scale the positions to the desired portfolio volatility.** Choose one positive multiplier that converts the raw sizes into dollar positions with the portfolio volatility you want. Calculate that volatility using the portfolio risk model, including correlations across positions.
6. **Adjust the dollar positions using the RD optimizer.** Control factor exposures and apply the portfolio's constraints. Recheck the resulting portfolio volatility after optimization.
7. **Perform sanity checks.**
 - a. **Concentration:** Is any position too large? Can the portfolio withstand a single-position crash, or a sharp squeeze in a short, under a realistic stress scenario?
 - b. **Related stocks:** Are there unintended bets in clusters of related companies? Reduce sizes where a common exposure makes the combined bet too large.
 - c. **Liquidity:** Are the positions sufficiently liquid to enter, adjust, and exit at the proposed sizes, including under stress?

Recheck portfolio risk after any further size adjustments.